

AutoDocx

A Unified 9-Agent Platform for Document Intelligence & Business Process Automation

9 specialised AI agents. One shared brain. Zero manual busywork.

Project & Business Plan Report

Prepared for the Student Innovation Grant (SIG)

Submitted by

Hardik Bisht | Prajyot Nula

Indian Institute of Science Education and Research (IISER), Bhopal

1. Objective of the Project

India's Micro, Small and Medium Enterprises (MSME) sector is the backbone of the national economy, comprising over 8.9 crore Udyam-registered enterprises and contributing close to 30% of India's GDP. Despite this scale, the overwhelming majority of these businesses continue to run their day-to-day operations on WhatsApp, spreadsheets, and paper-based records rather than on purpose-built digital infrastructure. Nine core business workflows that every growing enterprise depends on: document search and retrieval, employee onboarding, social media management, customer messaging, competitor and subsidy tracking, shipment monitoring, vendor evaluation, meeting-to-action conversion, and inbox management — remain almost entirely manual.

This manual dependence carries direct and measurable costs. Business-critical data is trapped across PDFs, scans, JPGs and presentations with no unified way to search or extract it. New hires wait days for answers to routine onboarding questions because HR teams must repeatedly field the same queries. WhatsApp, India's dominant business channel, still requires manual replies and order processing. Meeting notes rarely translate into tracked action items, so decisions get lost. Inboxes overflow, response service-level agreements slip, and customer trust erodes as a result.

The objective of this project AutoDocx is to close this automation gap through a single, unified platform: a shared multimodal Retrieval-Augmented Generation (RAG) core that powers nine specialised AI agents, each addressing one of the business-critical workflows above. Rather than requiring an SME to purchase, learn, and integrate nine separate point solutions, AutoDocx ingests a business's documents once PDFs, images, scanned invoices, presentations and handwritten records and reuses that shared understanding across every agent, from document search and HR onboarding to WhatsApp order automation, shipment tracking and meeting-note conversion.

The core value proposition can be summarised as: 9 specialised AI agents, one shared brain, zero manual busywork. Where existing solutions enterprise RPA platforms, document-AI specialists, or WhatsApp automation vendors each solve one narrow slice of the problem, AutoDocx is designed from the outset as an India-first, SME-first platform that bundles all nine workflows on a common data layer, making the product both cheaper to build (since every new agent reuses the same retrieval core) and cheaper to adopt (since customers pay only for the agents they need).

Through the Student Innovation Grant period, the project aims to move from an already-validated concept to a working shared RAG core and a first set of production-ready agents, laying the technical and commercial foundation for a startup that can be incorporated immediately after the grant-funded proof of concept.

2. Business Opportunity

Target Customers

AutoDocx's primary target customers are Indian SMEs and MSMEs employing between 10 and 500 people businesses that have outgrown ad hoc WhatsApp and Excel operations and generate enough document, communication and workflow volume to need automation, yet remain too small to afford enterprise RPA vendors or build in house AI teams. Within this base, the go-to-market plan prioritises three beachhead verticals: D2C/retail brands, manufacturing exporters, and professional services firms chosen for their high documentation loads, heavy WhatsApp driven customer interaction, and recurring onboarding and vendor-management needs.

Timeframe

The business plan targets incorporation of the startup within Month 6 of the grant period, immediately following the SIG-funded proof of concept, positioning the venture to begin commercial customer acquisition in its beachhead verticals as soon as the MVP is market-ready.

3. Solution Proposed

Parsing PPTX files and invoices through OCR and layout-aware parsers. This feeds a Shared Multimodal RAG Core the platform's central innovation which combines OCR and layout parsing, a vision-language model, document chunking, embeddings, a vector database, and retrieval-augmented LLM reasoning with guardrails, converting unstructured business documents into accurate, structured, queryable knowledge. An Agent Orchestration layer routes each incoming request to the correct specialised agent and shares context and memory across agents. A Channels layer exposes the platform through a web dashboard, WhatsApp, email, Slack/Teams and mobile AutoDocx is architected as a four-layer platform. The Ingestion layer accepts PDFs, JPGs/scanned images.

On top of this shared core sit nine specialised agents, summarised below, each reusing the same retrieval and reasoning infrastructure rather than being built as an isolated tool:

#	Agent	Core Function	Key Technology
1	Multimodal RAG	Search and extract across PDF, JPG, PPTX and scanned documents	OCR + VLM, embeddings, vector DB
2	Employee Onboarding	Answers new-hire questions from HR policy documents	RAG over policy docs, chat UI
3	Social Media Content	Drafts on-brand posts and captions on a content calendar	Brand-tuned LLM prompts, scheduling
4	WhatsApp Automation	Auto-replies, catalogue browsing, order flows	WhatsApp Business API, RAG
5	Capital Edge	Competitor pricing tracking and subsidy scheme matching	Web monitoring, scheme database
6	Shipment Tracker	Real-time shipment status and alerts	Carrier API integrations
7	Vendor Performance	Vendor scorecards and risk flags	Data aggregation, scoring model
8	Meeting Notes → Action	Converts meeting notes into tracked action items	Speech-to-text, summarisation
9	Email Helper	Prioritised inbox and auto-drafted replies	Mailbox integration, triage agent

The suggested technology stack a Claude/GPT-class LLM API, a pgvector/Milvus vector database, Tesseract/Azure Document Intelligence for OCR, LangGraph for orchestration, the WhatsApp Business API, and AWS/GCP cloud infrastructure relies entirely on mature, production-proven components, making the integration effort the primary engineering challenge rather than unproven research.

4. Novelty / Uniqueness

AutoDocx's core novelty lies in being the first unified suite to cover all nine SME workflows on a single, shared multimodal RAG core, rather than as nine disconnected point tools. Every existing alternative in the

market enterprise RPA platforms, document-AI specialists, WhatsApp automation vendors, or generic chatbot wrappers solves one narrow slice of the SME automation problem and requires a business to stitch multiple vendors together. AutoDocx instead ingests a business's documents once and reuses that shared understanding across every downstream agent, so each additional agent becomes progressively cheaper and faster to build than the last, since it inherits the retrieval, embedding and reasoning infrastructure already built for the core.

The platform is also explicitly designed India-first: it is WhatsApp-native (India's dominant business communication channel), supports multilingual OCR for the range of scripts used in Indian business documents, and is aware of GST filing and MSME-subsidy schemes relevant to Indian SMEs none of which are addressed by the primarily enterprise-oriented or English-first tools available today.

On intellectual property, while the platform's foundation models, OCR engines and vector-search components rely on mature, publicly available technology, the differentiated and protectable IP lies in the orchestration layer that routes and shares context between agents, the subsidy-matching dataset built specifically for Indian government schemes, and the WhatsApp automation flow templates developed for SME use cases. Patent filing for the orchestration approach is planned and in the interim, the prompt-engineering and orchestration pipelines will be protected as trade secrets, with access restricted within the founding team until formal IP protection is in place.

This combination architectural novelty in the “one core, many agents” design, and market novelty in India-first workflow coverage is what differentiates AutoDocx from any single function competitor, and is expected to compound into a defensible moat as more agents and more customer data are added to the shared core.

5. Customer Segment

The primary customer segment is Indian SMEs and MSMEs with 10 to 500 employees, a segment chosen deliberately because it sits between two poles poorly served by the current market. On one hand, enterprise automation vendors such as UiPath and Automation Anywhere price and design their offerings for large organisations with dedicated IT and RPA teams, well beyond the reach of a 50-person exporter or a 200-person D2C brand. On the other side, businesses below this threshold typically lack the transaction and document volume to justify any automation spend at all. The 10–500 employee band is the point at which document, communication and onboarding volumes become high enough that manual handling creates a real, recurring cost, while the business is still too small to build or buy enterprise-grade tooling.

Within this band, the go-to-market plan prioritises three beachhead verticals, to be narrowed to one or two during early customer validation: D2C/retail brands, which generate high WhatsApp order and customer-service volume; manufacturing exporters, who handle heavy invoice, shipment and vendor documentation; and professional services firms, which rely on document-heavy client work and frequent internal onboarding.

This segmentation is justified on three grounds. First, all three verticals sit squarely within India's 8.9 crore-strong Udyam-registered MSME base, so the segment is large and accessible. Second, each vertical independently experiences several of AutoDocx's nine target workflows at once for example, an exporter needs shipment tracking, vendor scoring and document extraction simultaneously making the “one shared brain” value proposition genuinely more compelling than any single point solution. Third, formal customer validation is already under way, with structured outreach to prospective SME customers and mentor feedback

forming part of the pre-incorporation plan, ensuring the segment choice is grounded in direct market feedback rather than assumption alone.

6. Market Analysis / Competitive Landscape

Market Size and Growth

The timing for AutoDocx is supported by three converging trends: rapid growth in the underlying agentic-AI and document-intelligence markets, a large and expanding MSME base, and a rising digital-readiness baseline in India (UPI, GST e-invoicing and Digital India initiatives) that makes SMEs increasingly comfortable adopting cloud-based tools.

Market	2026	2034	CAGR
Global Agentic AI (TAM)	~US\$9.1B	US\$139B	~40.5%
Global IDP (Adjacent)	~US\$14.2B	US\$91B	~26.2%
India MSME base	8.9 Cr+ registered	Growing ~15%/yr	
SAM – India SME automation spend	US\$12.8B	US\$18B	38.3%
SOM – Year 1–3 capture (target)	US\$450M	US\$1.5B	38.3%

Sources: Fortune Business Insights (2026); Ministry of MSME Udyam registration data.

Competitive Landscape

No existing point-solution bundles all nine core SME workflows on one shared, India-first data layer. The table below summarises the main categories of competitors and the gap each leaves open for AutoDocx.

Competitor Type	Examples	What They Cover	Gap They Leave Open
Enterprise RPA	UiPath, Automation Anywhere	Large-enterprise automation	Expensive, not SME-friendly
Point CRM/Helpdesk	Zoho, Freshworks	Sales/support workflows	No unified multimodal RAG
Document-AI specialists	ABBYY, Rossum, Hyperscience	Invoice/document extraction only	Single-function only
WhatsApp automation vendors	Gupshup, Interakt	WhatsApp messaging only	No shared data layer
Generic AI chatbot wrappers	Various	Basic Q&A over a knowledge base	No workflow depth

Because every existing alternative addresses only one workflow, an SME that wants document search, WhatsApp automation and vendor tracking today must buy and integrate three or more separate tools each with its own data silo. AutoDocx's shared data layer and modular, pay-for-what-you-use pricing directly target this gap: customers will opt for AutoDocx over piecemeal point solutions because it is the only offering that runs a business's core workflows on one shared brain, at SME-appropriate pricing and deployment speed.

7. Milestones / Scale-Up Strategy

The scale-up strategy is deliberately modular: rather than building all nine agents simultaneously, the platform is launched with two to three agents first, validated with real users, and then expanded incrementally, an approach made possible because every new agent reuses the already-built shared RAG core. Within the six-month grant period this translates into three sequential phases, shown below as a graphical work-plan.

Phase	Milestone	Target Duration	Key Deliverables
1	Working Prototype	3 months (M1–M3)	Shared multimodal RAG core + first 2 agents functional
2	Minimum Viable Product	1.5 months (M3.5–M4.5)	5–6 agents live, WhatsApp/Email channels, web dashboard
3	Market Launch	1.5 months (M4.5–M6)	All 9 agents integrated, billing and support systems live

Beyond the grant period, the scale-up strategy follows a yearly cadence. In Year 1, the focus is on converting the pilot cohort into the first paying customers within the beachhead verticals and reaching a steady cadence of monthly agent releases. In Year 2, the plan is to expand into adjacent verticals and geographies, introduce the enterprise/white-label tier, and deepen the subsidy-matching and multilingual OCR capabilities that anchor the platform's India-first positioning. In Year 3, the objective is to scale customer acquisition toward the targeted US\$450 million Serviceable Obtainable Market, supported by the compounding cost advantage of the shared-core architecture.

In the future, we can also take help from Indian government schemes like Startup India.

8. Project Timelines

The Student Innovation Grant period runs for six months. In line with grant requirements, the project defines clear monthly milestones so that progress can be tracked against plan by the Expert Panel/Selection Committee. The table below sets out the month-by-month work plan for the grant period.

Month	Planned Activity
M0	Requirements gathering and system architecture design
M1	Build the shared multimodal RAG core (ingestion, OCR, embeddings, vector DB, retrieval)
M2	Build Agents 1–3: Multimodal RAG, Employee Onboarding, Email Helper
M3	Build Agents 4–6: WhatsApp Automation, Social Media Content, Meeting Notes → Action
M4	Build Agents 7–9: Capital Edge, Shipment Tracker, Vendor Performance
M5	Integration testing, pilot deployment with real users, and MVP freeze

Each monthly milestone is paired with a Technology Readiness Level (TRL) target, giving the selection committee an objective way to track technical maturity across the grant period, from initial core validation (TRL 4) through to a market-ready MVP (TRL 7):

Duration	Milestone Target	TRL (1–9)
Months 1–2	Shared multimodal RAG core functional across PDF, JPG, PPTX and scanned documents	3
Months 2–3	First 3 agents live in a test environment	2
Months 3–4	WhatsApp Automation and Social Content agents integrated	1
Months 4–5	All 9 agents integrated; pilot launched with real users	1
Months 5–6	Pilot feedback incorporated; MVP ready for market	1

9. Expected Outcome

By the end of the six-month grant duration, AutoDocxff expects to deliver a working Minimum Viable Product with all nine agents integrated on the shared multimodal RAG core, accessible through a web dashboard and WhatsApp/email channels, with basic billing and support functionality in place. Technology Readiness Level is expected to progress from an initial TRL of 4 core retrieval technology validated in a relevant environment to TRL 7, at which the system is demonstrated as a market-ready prototype validated through a real-user pilot.

Specific, measurable success criteria to be finalised and tracked over the grant period include: document extraction accuracy across the five supported document types (PDF, JPG/scanned image, PPTX, invoices and handwritten documents); end-to-end response latency per query; and agent task-completion accuracy measured on a test set of real business documents and queries drawn from pilot customers.

Beyond the technical deliverable, the project expects to generate measurable business impact for its pilot users: an estimated 15–20+ hours per week of manual busywork automated per business; materially faster new-employee onboarding and customer response times; and fewer missed subsidy or competitor-intelligence opportunities and document-processing errors, as document intelligence shifts from manual, siloed effort to a shared, automated layer.

On the commercial side, the expected outcome of the grant period is a validated, working shared-core architecture that de-risks the remaining product roadmap since building agent ten onward reuses proven infrastructure rather than requiring nine separate research bets together with structured pilot feedback from real SME users. This is expected to directly support incorporation of the startup within Month 6, immediately following the SIG-funded proof of concept, and a clear, fast path from PoC to pilot to early revenue.

